

How to Implement the Lambda Calculus, Quickly

- Stephanie Weirich
- University of Pennsylvania
- Haskell Love
- September 10, 2021

<https://github.com/sweirich/lambda-n-ways/>

Key lambda-calculus operation

Capture-avoiding substitution

$a \{ b / x \}$

$(z (\lambda y. z)) \{ y / z \}$

\Rightarrow

$y (\lambda w. y)$



Klaus Ostermann

@klaus03



What I tell others what I do: Design cool programming languages. What I really do: Debug substitution functions.

6:08 AM · Sep 3, 2021 · Twitter Web App

8 Retweets **1** Quote Tweet **76** Likes

Magic implementation?

```
data Exp = Var Var          -- x
         | Lam (Bind Exp)    --  $\lambda x.a$ 
         | App Exp Exp       -- (a b)
         deriving (Substitution)

-- definition of substitution "for free"
-- subst :: Var -> Exp -> Exp -> Exp  -- a { b / x }
```

Binding Library (unbound-generics)

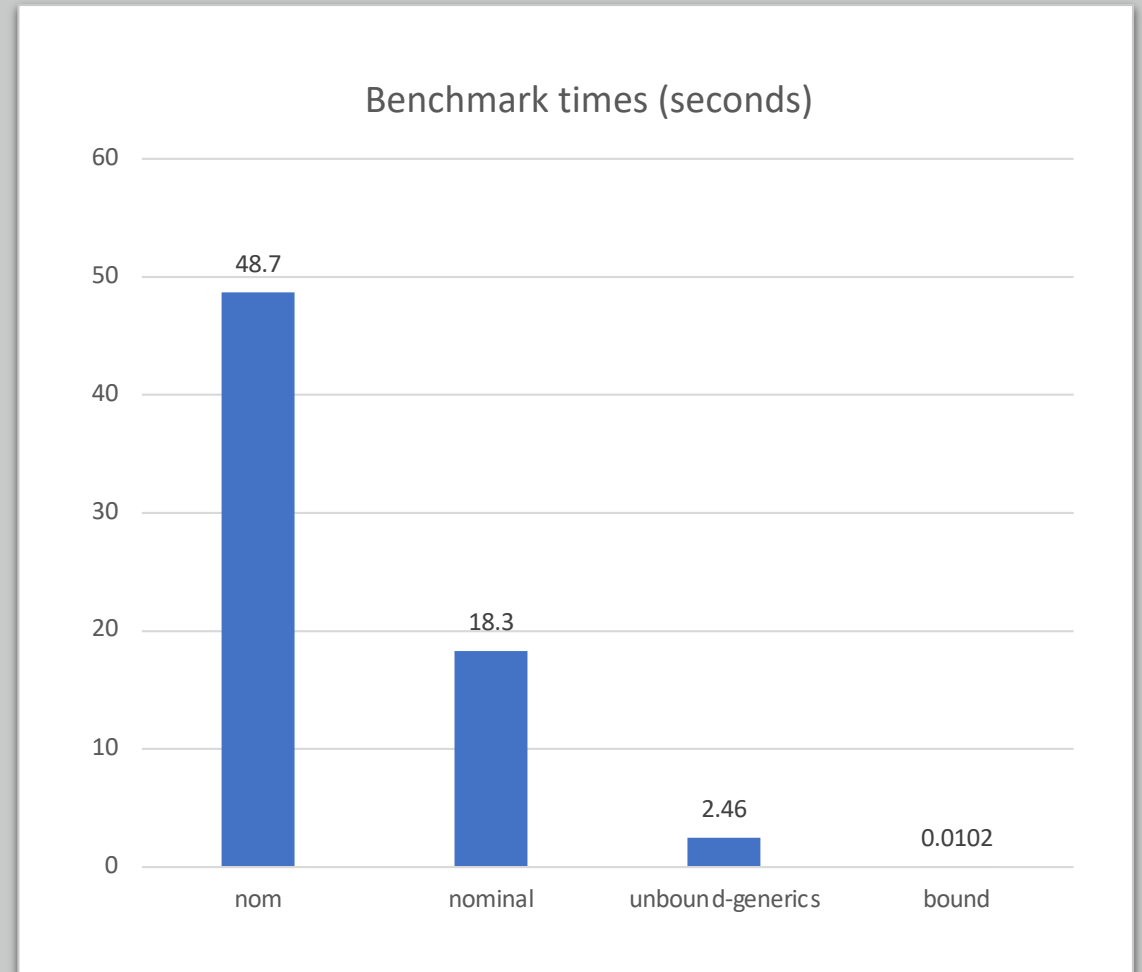
```
import LocallyNameless.UnboundGenerics

data Exp = Var (Name Exp)           -- x
        | Lam (Bind (Name Exp) Exp) --  $\lambda x. a$ 
        | App Exp Exp               -- (a b)
        deriving (Show, Generic)

instance Subst Exp Exp where
  isvar (Var x) = Just (SubstName x)
  isvar _       = Nothing
  -- quick definition of substitution
```

Multiple libraries available on hackage!

- Unbound (GHC version $\leq 8.8.3$)
- unbound-generics
- bound
- nominal
- nom



GHC 8.8.3, MacBook pro, 2.4 GHz 8-Core Intel Core i9, 64 GB, measured using criterion

Why the difference?

- Different developers?
 - Reimplement uniformly
- Extra features? (pattern binding, annotations, etc.)
 - Only reimplement "core" mechanism
- Magic GHC pragma? `{-# SUMMON SPJ #-}`
 - Study GHC manual
- Different generic programming?
 - Compare with and without
- Different optimizations?
 - Cross-fertilize ideas
- Different binding representations?

- **nom/nominal**
Name-based binding
- **bound**
de Bruijn indices
- **Unbound/unbound-generics**
Locally nameless (mixed)

This work

- Start with three basic implementations
 - Names and simple renaming
 - de Bruijn indices
 - Locally nameless representation
- Create optimized versions
- Use type classes and generics to define a library interface
- Benchmark

Inspiration

λ -calculus cooked four ways

Lennart Augustsson

1 Introduction

This little paper describes how to implement λ -calculus in four different ways. To be precise, it shows how to implement the functions that compute the (β) normal form of an expression.

2005-2006, <https://github.com/steshaw/lennart-lambda>

Lennart's benchmark

Normal-order reduction of the Church encoding of

$6! == \text{sum } [1 \dots 37] + 17$

i.e. $720 == 719$

Benchmark statistics

- Number of substitutions required for normalization: **119,697**
- AST depth: 53, binding depth: 25
- Average # of variable occurrences during each beta-reduction: 1.15

Example: Computing normal form with names

$\lambda w. \lambda x. (\lambda y. w y (\lambda x. w y (\lambda y. w y))) (x x x)$
substitute $(x x x)$ for y

$\lambda w. \lambda x. w (x x x) (\lambda z. w (x x x) (\lambda y. w y))$
rename x to avoid capture stop when y not free

Example: with de Bruijn indices

$\lambda w. \lambda x. (\lambda y. w y (\lambda x. w y (\lambda y. w y))) (x x x)$

$\lambda. \lambda. (\lambda. 2\ 0 (\lambda. 3\ 1 (\lambda. 4\ 0))) (0\ 0\ 0)$

substitute (0 0 0) for 0

$\lambda. \lambda. \underline{1} (0\ 0\ 0) (\lambda. \underline{2} (\underline{1\ 1\ 1}) (\lambda. \underline{3}\ 0))$

decrement other variables increment (0 0 0) under binders
after reduction

Example: with Locally nameless

$\lambda w. \lambda x. (\lambda y. w y (\lambda x. w y (\lambda y. w y))) (x x x)$

name outermost variables to expose beta-reduction

$\lambda. \lambda. (\lambda. 2 \ 0 (\lambda. 3 \ 1 (\lambda. 4 \ 0))) (0 \ 0 \ 0)$

$(\lambda. w \ 0 (\lambda. w \ 1 (\lambda. w \ 0))) (x \ x \ x)$

substitute (x x x) for 0

replace names with indices when finished

$w (x \ x \ x) (\lambda. w (x \ x \ x) (\lambda. w \ 0))$

$\lambda. \lambda. 1 (0 \ 0 \ 0) (\lambda. 2 (1 \ 1 \ 1) (\lambda. 3 \ 0))$

Bookkeeping during $b \{a/x\}$

- Names: **rename bound variables to avoid capture**
 - calculate free variables of a
 - rename bound variables in b with a "fresh" name
- de Bruijn: **adjust indices**
 - shift free indices in a depending on binding depth
 - decrement indices of b because we lost a binder
- Locally nameless: **exchange names/indices**
 - invariant: a has no "free" indices, only free names
 - exchange happens during traversal, not substitution
 - need to choose "fresh" names

```
data Exp =  
    Var Var  
  | Lam (Bind Exp)  
  | App Exp Exp
```

What does it look like to use
these three approaches?

Let's implement normal-order reduction

Normal-order full reduction w/ names

```
nf :: Exp -> Exp
nf (Var x) = Var x
nf (Lam (Bind x e)) =
    -- recurse under binder
    Lam (Bind x (nf e))
nf (App a b) =
    case whnf a of
        Lam (Bind x a0) ->
            nf (subst x b a0)
        a' -> App (nf a') (nf b)
```

```
whnf :: Exp -> Exp
whnf (Var x) = Var x
whnf (Lam b) =
    -- don't recurse
    Lam b
whnf (App a b) =
    case whnf a of
        Lam (Bind x a0) ->
            whnf (subst x b a0)
        a' -> App a' b
    -- normalize head only
```


Normal-order full reduction w/ indices

```
nf :: Exp -> Exp
nf (Var x) = Var x           -- x is an Int (index)
nf (Lam (Bind e)) =          -- no Var stored at binder
    -- recurse under binder
    Lam (Bind (nf e))
nf (App a b) =
    case whnf a of
        Lam (Bind a0) ->
            nf (subst 0 b a0) -- shift indices during subst
        a' -> App (nf a') (nf b)
```

Reduction w/ locally nameless terms

```
nf :: Exp -> Exp
nf (Var x) = (Var x)           -- invariant, no free indices
nf (Lam (Bind b)) =            -- no Var at binder
  let x = fresh b              -- find var not free in b
      b' = nf (subst 0 (Var x) b) -- cvt index to name
  in Lam (Bind (close x b'))    -- cvt name to index
nf (App a b) =
  case whnf a of
    Lam (Bind a0) ->
      nf (subst 0 b a0)         -- no shifting (from invariant)
    a' -> App (nf a') (nf b)
```

But which is faster?

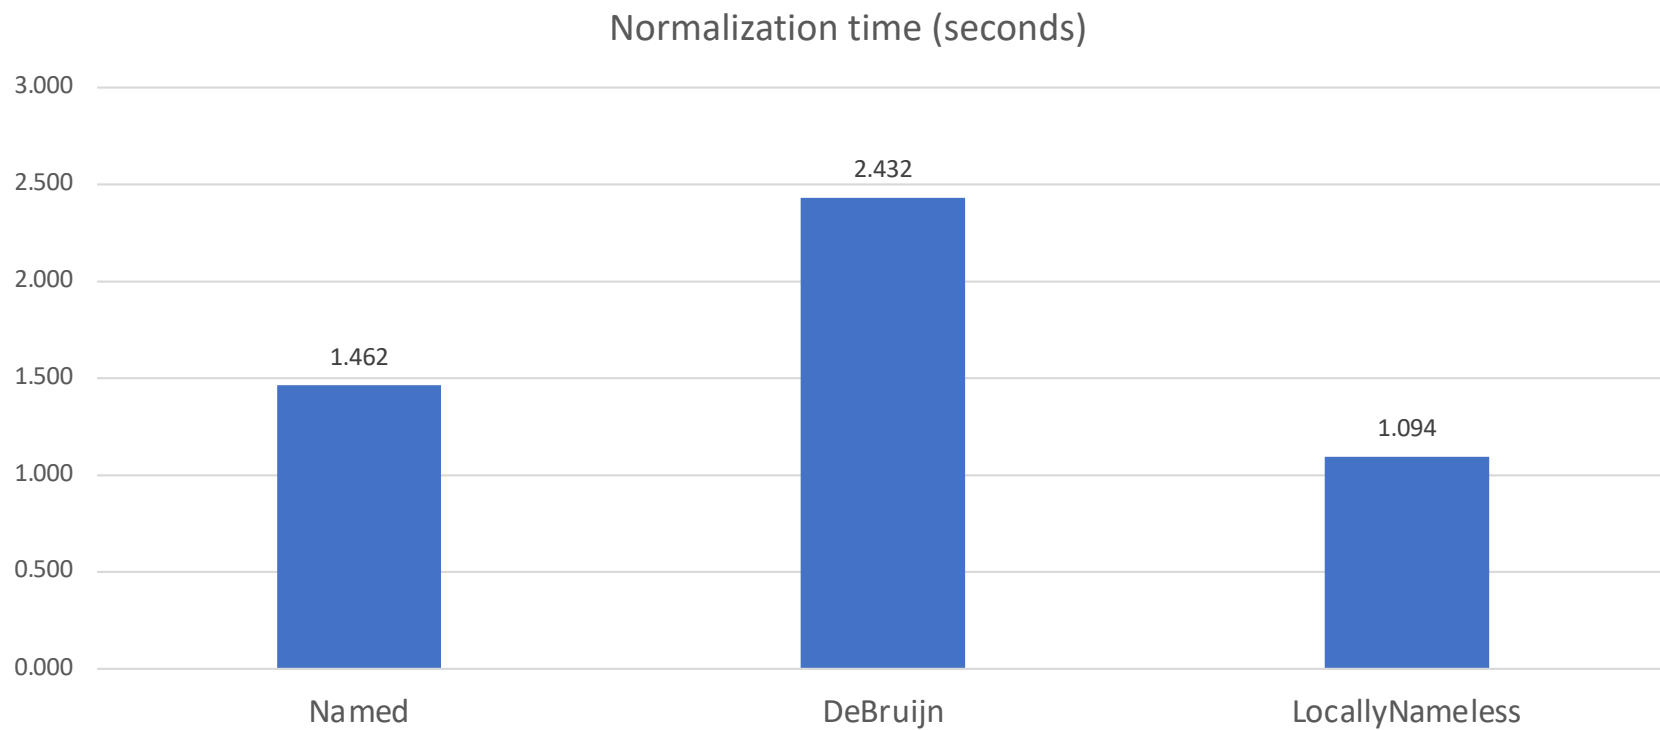
<https://github.com/sweirich/lambda-n-ways>

Lennart.Simple – adapted from "Lambda-calculus cooked four ways", with bugfix

Lennart.DeBruijn – adapted from "Lambda-calculus cooked four ways"

LocallyNameless.Ott – generated by Ott tool

Benchmark results: basic versions head-to-head



Optimize!

Named.Lazy.SimpleH

DeBruijn.Lazy.Par.B

LocallyNameless.Lazy.Opt



Yaron (Ron) Minsky

@yminsky



Here's a nice PR to OCaml that's fixing some performance bugs in the typechecker.

ocaml/ocaml

#10599 Evaluate signature substitutions lazily



2 comments 5 reviews 16 files +643 -287



stedolan • September 1, 2021 10 commits



Evaluate signature substitutions lazily by stedolan · Pull Request #10599 · oca...
When typechecking code that imports large libraries, the compilers spends much of its time evaluating substitutions in signatures (freshening identifiers, renami...
[github.com](#)

6:27 AM · Sep 7, 2021 · Twitter Web App

1 Retweet 22 Likes

Can we do better? Yes!

- In all three versions, bookkeeping leads to multiple traversals
- General Idea: generalize traversals and cache info at binders

`subst :: Sub Exp -> Exp -> Exp`

- **For Names**

- Multi-substitution `type Sub a = Map Var a`
- Cache free vars `data Bind a = Bind VarSet Var a`

- **Benefits**

- Find fresh variables quickly
- Can cut off substitution early if domain doesn't affect free variables
- Fewer traversals: fuse renaming and normal substitutions

Can we do better? Yes!

- In all three versions, bookkeeping leads to multiple traversals
- General Idea: generalize traversals and cache info at binders

`subst :: Sub Exp -> Exp -> Exp`

- **For de Bruijn indices**

- Multi-substitution `data Sub a = Inc Int | Cons a (Sub a) | Compose (Sub a) (Sub a)`
- Delay substitution `data Bind a = Bind (Sub a) Var a`

- **Benefits**

- Compose substitutions using smart constructors
- Fewer traversals: multiple indices replaced simultaneously

Can we do better? Yes!

- In all three versions, bookkeeping leads to multiple traversals
- General Idea: generalize traversals and cache info at binders

`subst :: Sub Exp -> Exp -> Exp`


- **For Locally Nameless**

- Multi-subst/close `type Sub a = (Int, [Exp])`
- Delay last traversal `data Bind a = Bind (Info a) Var a`
 `data Info a = NoInfo | Subst Int [a]`
 `| Close Int [IdInt]`

- **Benefits**

- Fewer traversals: multiple indices/names replaced simultaneously

Comparison: Original vs. optimized



```
data Exp =  
    Var Var  
  | Lam (Bind Exp)  
  | App Exp Exp
```

Make it generic!

Named.Lazy.SimpleGH

DeBruijn.Lazy.Par.GB

LocallyNameless.Lazy.GenericInstOpt

Named Client: virtually no code

```
instance VarC Exp where
```

```
  var = Var
```

```
  isvar (Var v) = Just v
```

```
  isvar _ = Nothing
```

```
instance FreeVarsC Exp where
```

```
instance SubstC Exp Exp where
```

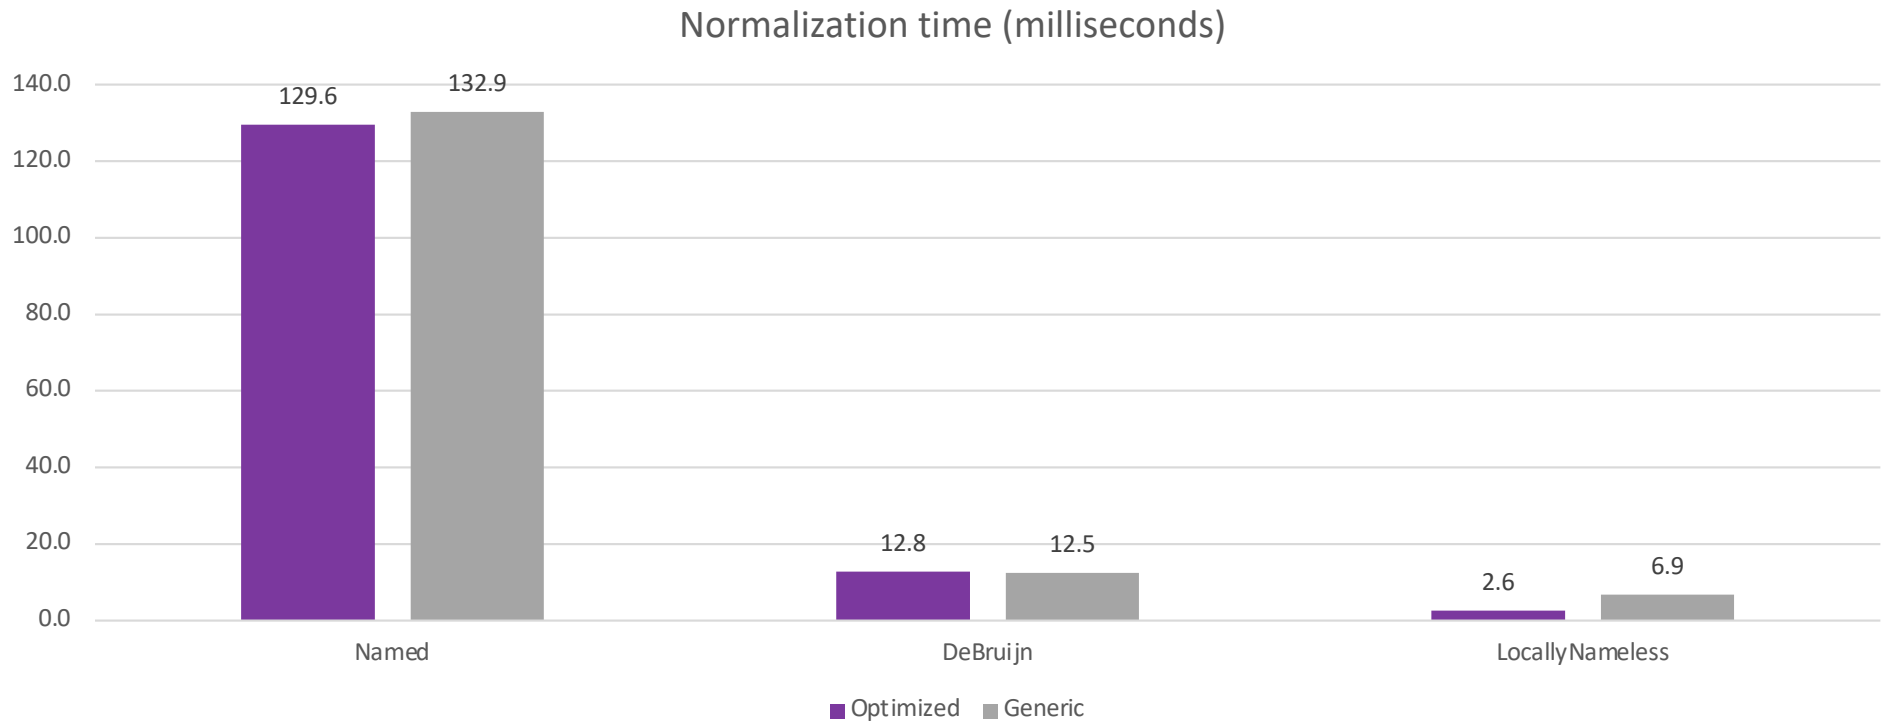
Named Client: virtually no code

```
instance VarC Exp where
  {-# SPECIALIZE instance VarC Exp #-}
  var = Var
  isvar (Var v) = Just v
  isvar _ = Nothing
```

```
instance FreeVarsC Exp where
  {-# SPECIALIZE instance FreeVarsC Exp #-}
```

```
instance SubstC Exp Exp where
  {-# SPECIALIZE instance SubstC Exp Exp #-}
```

Benchmark: Optimized vs. Generic



```
data Exp =
```

```
    Var Var
```

```
  | Lam (Bind Exp)
```

```
  | App Exp Exp
```

```
data Exp =
```


```
    Var {-# UNPACK #-} !Var
```

```
  | Lam !(Bind Exp)
```

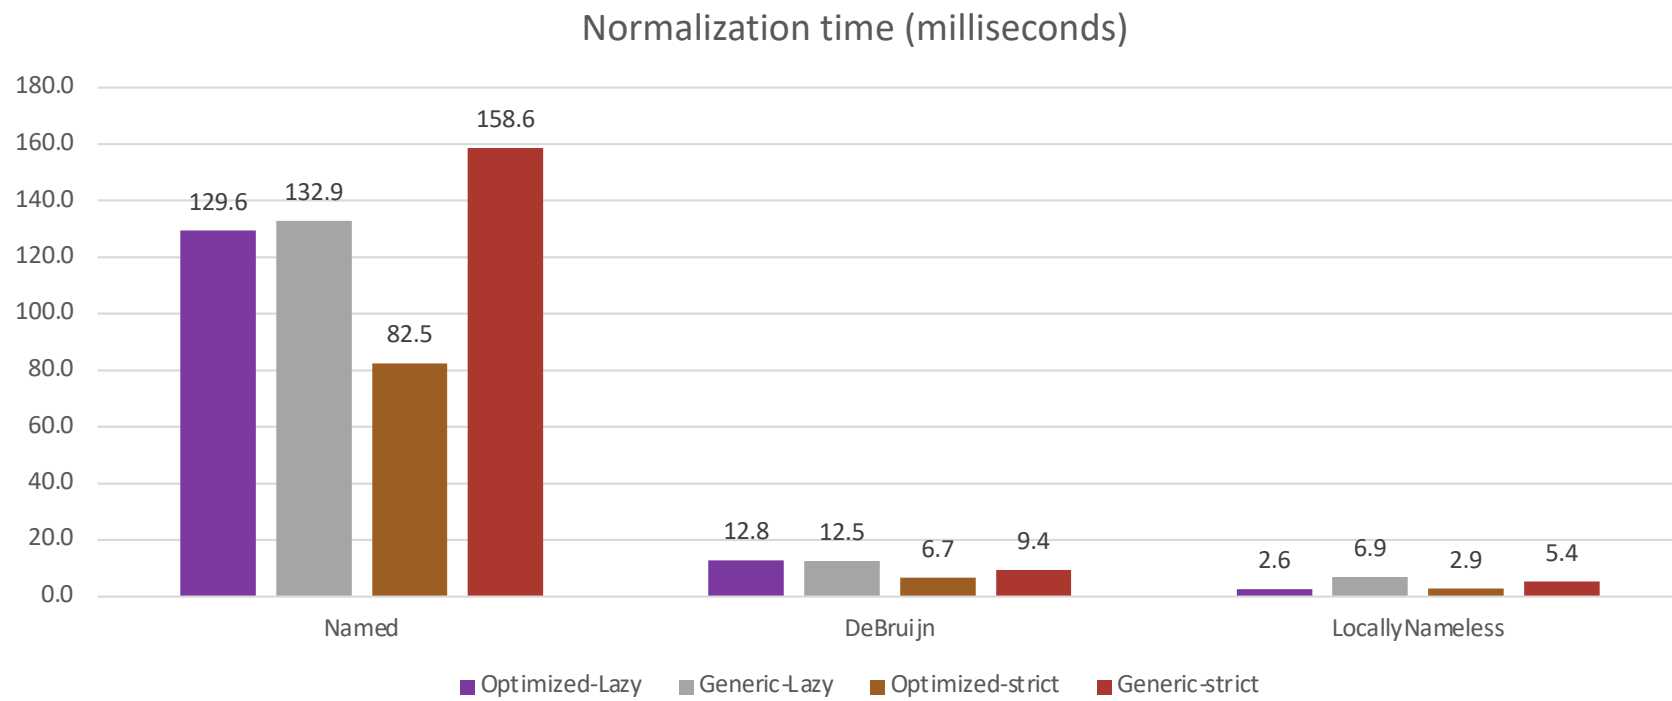
```
  | App !Exp !Exp
```

Coda: strictness?

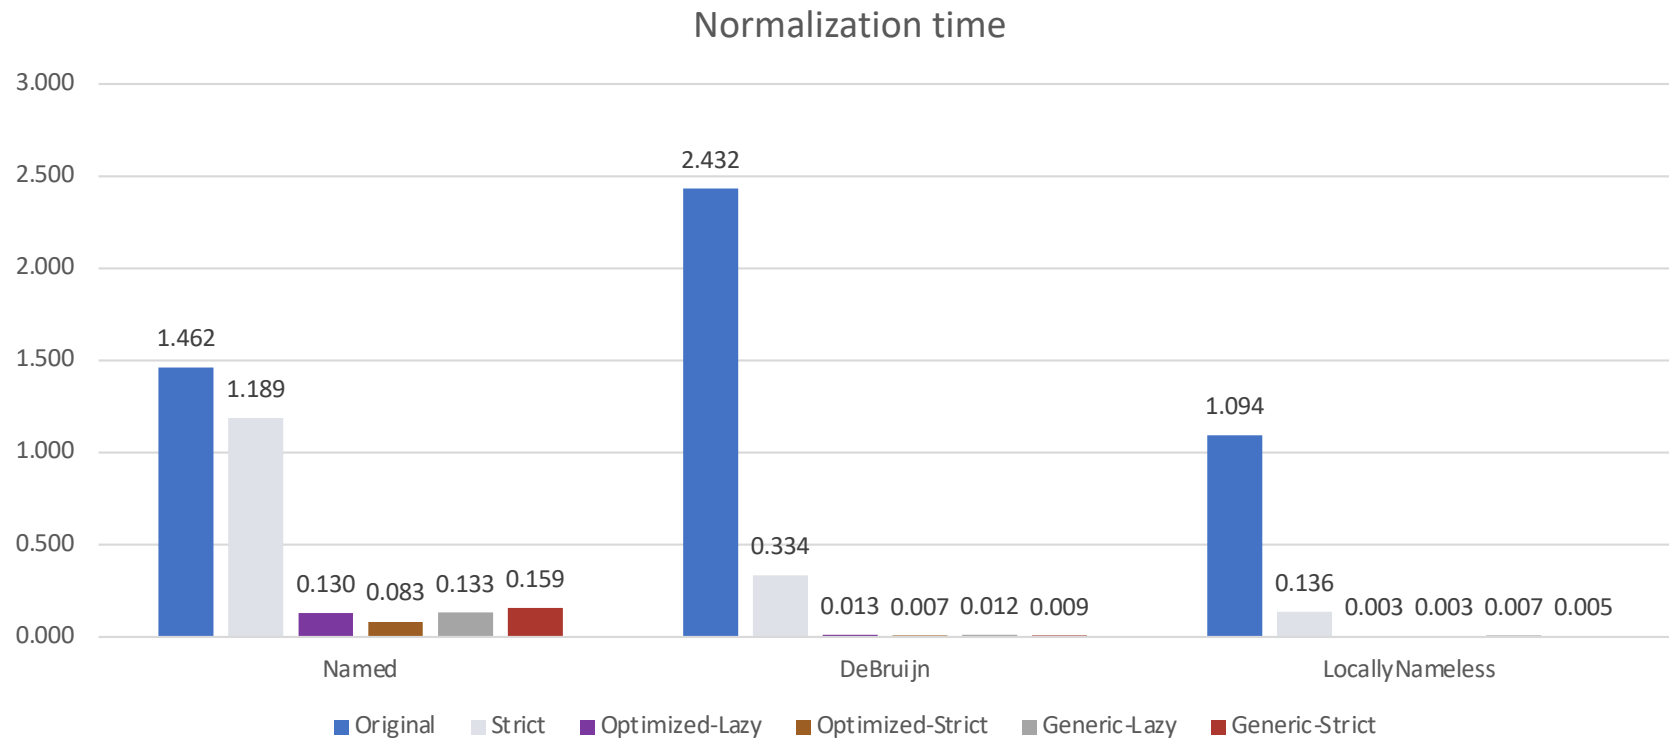
Benchmark results: strictness annotations



Strictness annotations and optimization



Benchmark summary



Conclusions

- This is **one** benchmark, so don't read *too* much into it, but
 - Optimization more significant than representation
 - Generic programming is relatively low runtime cost
- More to do even, in this context
 - de Bruijn optimizations not the same as "bound"
 - locally nameless optimizations not the same as "unbound"
 - Many other alternatives
- Contributions welcome!
<https://github.com/sweirich/lambda-n-ways>